

Homework 6 Solutions - MATH 4322

Cathy Poliak

Spring 2024

Problem 1

You are given:

- n data samples $\mathbf{x}_i = (x_{1,i}, \dots, x_{p,i}), i = 1, \dots, n$
- n corresponding to true responses (or labels) $y_i, i = 1, \dots, n$.

and asked to train a single linear neuron “network” to approximate function $f(\cdot)$ such that $f(\mathbf{x}_i) = y_i, i = 1, \dots, n$. Provided the train steps for your “network” by answering the following questions.

- a) What is the formula to calculate an output \hat{y}_i from an input \mathbf{x}_i ? What are the model parameters in that formula?
- b) What criteria do we need to optimize in order to estimate the model parameters?
- c) What is the name of the method used to optimize this criteria in case you do not have access to an analytical solution?

Answers

- a) $\hat{y}_i = \sum_{i=1}^p w_i x_i + b$, the model parameters are:
 - Weights of inputs: w_1, w_2, \dots, w_p .
 - bias: b .
- b) We need to optimize the sum of the squared residuals of predicted values \hat{y}_i , when compared to the true values y_i :

$$\frac{1}{2} \sum_{i=1}^p (\hat{y}_i - y_i)^2$$

Or if your data samples are fed one at a time, then when the i^{th} sample is fed, we need to optimize:

$$\frac{1}{2} (\hat{y}_i - y_i)^2$$

- c) Gradient decent.

Problem 2

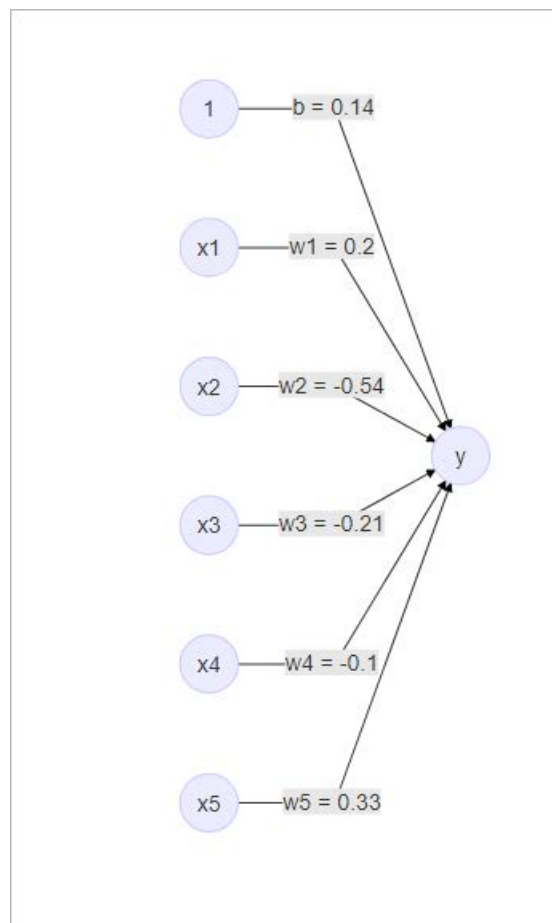
Presume that for a single linear neuron model with input variables x_1, \dots, x_5 , you are given the following parameter values:

- weights: $w_1 = 0.2, w_2 = -0.54, w_3 = -0.21, w_4 = -0.1, w_5 = 0.33$,
- bias: $b = 0.14$.

- Draw a mathematical model of this linear neuron that takes an arbitrary input vector $\mathbf{x} = (x_1, x_2, x_3, x_4, x_5)$.
- Calculate the linear neuron output for the case of $x_1 = 4, x_2 = -3, x_3 = 7, x_4 = 5, x_5 = -1$. Show your work.

Answer

- Diagram



b) $\hat{y} = 0.14 + 0.2(4) - 0.54(-3) - 0.21(7) - 0.1(5) + 0.33(-1) = 0.26$

Problem 3

You are given an artificial neural network (ANN) of linear neurons with

- Input layer of two neurons: x_1, x_2
- Fully-connected hidden layer of three neurons: h_1, h_2, h_3
- One output neuron, y .

The following weight matrices are provided:

1) Between input & hidden layer:

		Hidden h_1	h_2	h_3
	1 (bias)	-0.3	0.5	0.5
Input	x_1	0.6	-0.4	0.5
	x_2	-0.7	-0.3	0.2

2) Between hidden & output layer:

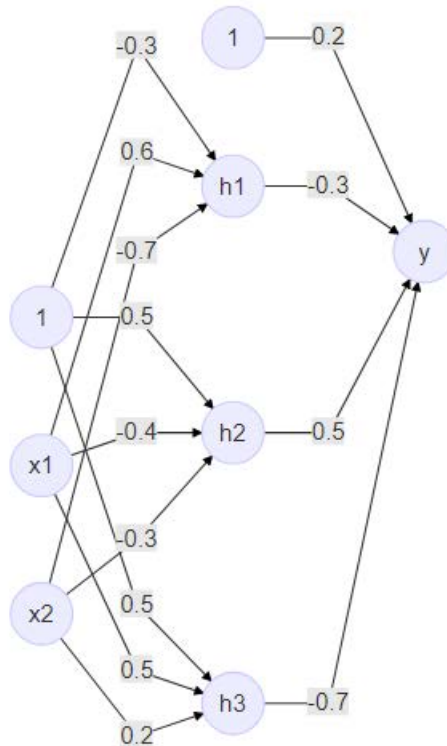
		Output y
	1 (bias)	0.2
Hidden	h_1	-0.3
	h_2	0.5
	h_3	-0.7

a) Draw this ANN as was done in lecture slides.

b) Calculate the output of this ANN for the case of $x_1 = 10, x_2 = -5$. Show work.

Answers

a) Diagram



b) Calculations:

$$\begin{aligned}\hat{h}_1 &= -0.3 + 0.6 * 10 - 0.7 * (-5) = 9.2 \\ \hat{h}_2 &= 0.5 - 0.4 * 10 - 0.3 * (-5) = -2 \\ \hat{h}_3 &= 0.5 + 0.5 * 10 + 0.2 * -5 = 4.5 \\ \hat{y} &= 0.2 - 0.3 * 9.2 + 0.5 * -2 - 0.7 * 4.5 = -6.71\end{aligned}$$

Problem 4

- We want to predict the ‘medv’ value based on the input of the other thirteen variables.
- We will run a regression neural network for the Boston data set.
- We will split the data into training/testing by a 70/30 split.

a) Type and run the following in R.

```
library(neuralnet)
library(MASS)

data = Boston #renaming the Boston data set to "data"
summary(data)
```

##	crim	zn	indus	chas
##	Min. : 0.00632	Min. : 0.00	Min. : 0.46	Min. : 0.00000
##	1st Qu.: 0.08205	1st Qu.: 0.00	1st Qu.: 5.19	1st Qu.: 0.00000
##	Median : 0.25651	Median : 0.00	Median : 9.69	Median : 0.00000
##	Mean : 3.61352	Mean : 11.36	Mean : 11.14	Mean : 0.06917
##	3rd Qu.: 3.67708	3rd Qu.: 12.50	3rd Qu.: 18.10	3rd Qu.: 0.00000
##	Max. : 88.97620	Max. : 100.00	Max. : 27.74	Max. : 1.00000
##	nox	rm	age	dis
##	Min. : 0.3850	Min. : 3.561	Min. : 2.90	Min. : 1.130
##	1st Qu.: 0.4490	1st Qu.: 5.886	1st Qu.: 45.02	1st Qu.: 2.100
##	Median : 0.5380	Median : 6.208	Median : 77.50	Median : 3.207
##	Mean : 0.5547	Mean : 6.285	Mean : 68.57	Mean : 3.795
##	3rd Qu.: 0.6240	3rd Qu.: 6.623	3rd Qu.: 94.08	3rd Qu.: 5.188
##	Max. : 0.8710	Max. : 8.780	Max. : 100.00	Max. : 12.127
##	rad	tax	ptratio	black
##	Min. : 1.000	Min. : 187.0	Min. : 12.60	Min. : 0.32
##	1st Qu.: 4.000	1st Qu.: 279.0	1st Qu.: 17.40	1st Qu.: 375.38
##	Median : 5.000	Median : 330.0	Median : 19.05	Median : 391.44
##	Mean : 9.549	Mean : 408.2	Mean : 18.46	Mean : 356.67
##	3rd Qu.: 24.000	3rd Qu.: 666.0	3rd Qu.: 20.20	3rd Qu.: 396.23
##	Max. : 24.000	Max. : 711.0	Max. : 22.00	Max. : 396.90
##	lstat	medv		
##	Min. : 1.73	Min. : 5.00		
##	1st Qu.: 6.95	1st Qu.: 17.02		
##	Median : 11.36	Median : 21.20		
##	Mean : 12.65	Mean : 22.53		
##	3rd Qu.: 16.95	3rd Qu.: 25.00		
##	Max. : 37.97	Max. : 50.00		

What is the mean of age? What is the mean of ptratio?

Answer

mean(age) = 68.57, mean(ptratio) = 19.05

b) Normalizing data

- It is recommended to **normalize** (or scale, or standardize, either works) features in order for all the variables to be on the same scale.
- With normalization, data units are eliminated, allowing you to easily compare data from different locations.
- This avoids unnecessary results or difficult training processes resulting in algorithm convergence problems.
- There are different methods for scaling the data.
- The **z-normalization**

$$x_{scale} = \frac{x - \bar{x}}{s}$$

- The **min-max scale**

$$x_{scale} = \frac{x - x_{min}}{x_{max} - x_{min}}$$

- And so forth
- The function in R is `scale(x, center = , scale =)`
- For this example we will use the min-max method to get all the scaled data in the range [0, 1].
- In order to scale we need to find the minimum and maximum value for each of the columns in the data set. To do this we use the `apply` function.
- The `apply` function returns a vector or an array or a list of values obtained by **applying** a function to margins of an array or matrix.
- Type and run the following:

```
max_data = apply(data,2,max)
#2=columns, we are getting the maximum value from each column
min_data = apply(data,2,min)
data_scaled = scale(data, center = min_data, scale = max_data - min_data)
head(data_scaled)
```

```
##          crim    zn      indus chas      nox      rm      age      dis
## 1 0.0000000000 0.18 0.06781525    0 0.3148148 0.5775053 0.6416066 0.2692031
## 2 0.0002359225 0.00 0.24230205    0 0.1728395 0.5479977 0.7826982 0.3489620
## 3 0.0002356977 0.00 0.24230205    0 0.1728395 0.6943859 0.5993821 0.3489620
## 4 0.0002927957 0.00 0.06304985    0 0.1502058 0.6585553 0.4418126 0.4485446
## 5 0.0007050701 0.00 0.06304985    0 0.1502058 0.6871048 0.5283213 0.4485446
## 6 0.0002644715 0.00 0.06304985    0 0.1502058 0.5497222 0.5746653 0.4485446
##          rad      tax  ptratio    black    lstat    medv
## 1 0.00000000 0.20801527 0.2872340 1.0000000 0.08967991 0.4222222
## 2 0.04347826 0.10496183 0.5531915 1.0000000 0.20447020 0.3688889
## 3 0.04347826 0.10496183 0.5531915 0.9897373 0.06346578 0.6600000
## 4 0.08695652 0.06679389 0.6489362 0.9942761 0.03338852 0.6311111
## 5 0.08695652 0.06679389 0.6489362 1.0000000 0.09933775 0.6933333
## 6 0.08695652 0.06679389 0.6489362 0.9929901 0.09602649 0.5266667
```

What is the scaled value of the first observation for medv?

Answer 0.42222

- c) Now we can split the data into training and testing data sets. We will use the 70/30 split

```
set.seed(10)
index = sample(1:nrow(data), round(0.7*nrow(data)))
train_data = as.data.frame(data_scaled[index,])
test_data = as.data.frame(data_scaled[-index,])
dim(train_data)
```

```
## [1] 354 14
```

How many observations do we have in the training data set?

Answer 354

d) Type and run the following

```
set.seed(1)
net_data = neuralnet(medv ~ ., data = train_data,
                      hidden = 10, linear.output = TRUE)

plot(net_data)
```

Apply the test data set to determine the MSE

```
predict_net = predict(net_data, test_data)
predict_net_start = predict_net * (max(data$medv) - min(data$medv)) + min(data$medv)
test_data_start = test_data$medv * (max(data$medv) - min(data$medv)) + min(data$medv)
sum((predict_net_start - test_data_start)^2) / nrow(test_data)
```

```
## [1] 15.37819
```

What is the test MSE for this model?

Answer I got 15.37819, this is \$3921.5.

e) Let us compare this test MSE to the linear regression model. Type and run the following:

```
lm.boston = lm(medv ~ ., data = data, subset = index)
summary(lm.boston)
```

```
##
## Call:
## lm(formula = medv ~ ., data = data, subset = index)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -13.981  -3.029  -0.529   1.882  25.349
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  4.073e+01  6.167e+00   6.606 1.53e-10 ***
## crim        -9.253e-02  5.445e-02  -1.700 0.090140 .
## zn           4.810e-02  1.737e-02   2.769 0.005936 **
## indus       -6.504e-04  8.068e-02  -0.008 0.993573
## chas         2.910e+00  1.124e+00   2.589 0.010045 *
## nox        -1.948e+01  4.636e+00  -4.202 3.39e-05 ***
## rm          3.203e+00  5.057e-01   6.333 7.60e-10 ***
## age         1.022e-02  1.627e-02   0.628 0.530150
## dis        -1.481e+00  2.454e-01  -6.033 4.20e-09 ***
## rad         2.948e-01  8.303e-02   3.551 0.000439 ***
## tax        -1.227e-02  4.786e-03  -2.564 0.010775 *
## ptratio     -9.461e-01  1.678e-01  -5.638 3.61e-08 ***
## black        9.767e-03  3.464e-03   2.819 0.005094 **
## lstat       -5.385e-01  6.229e-02  -8.646 < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
##
## Residual standard error: 4.989 on 340 degrees of freedom
## Multiple R-squared:  0.7062, Adjusted R-squared:  0.6949
## F-statistic: 62.85 on 13 and 340 DF,  p-value: < 2.2e-16

test = data[-index,]
predict_lm = predict(lm.boston,test)
sum((predict_lm - test$medv)^2)/nrow(test)

## [1] 17.77379
```

What is the training MSE for the linear model?

Answer I got 17.77379 which is higher than the neural network.