

# Bagging, Random Forest and Boosting

## Lab 12 - MATH 4322

- We will apply bagging, random forests and boosting to the `Boston` data, using the `randomForest` package.
- *Note:* The exact results obtained in this lab may depend on the version of R and the version of the `randomForest` package installed on your computer. Give the results from your computer.
- You can use the Rmarkdown script given or write down your answers and scan them as a pdf file to upload in Canvas similar to your homework.
- Possible points: 10.

**Question 1:** For any data that has  $p$  predictors **bagging** requires that we consider how many predictors at each split in a tree?

Bagging will use all  $p$  predictors in a tree

First, we call the data and create training/testing sets.

```
library(ISLR2)
set.seed(1)
train = sample(1:nrow(Boston), nrow(Boston)/2)
boston.test = Boston[-train, "medv"]
```

## Bagging

We perform bagging as follows:

```
library(randomForest)
set.seed(10)
bag.boston = randomForest(medv~., data = Boston,
                           subset = train,
                           mtry = ncol(Boston) - 1,
                           importance = TRUE)
```

```
bag.boston
```

Call:

```
randomForest(formula = medv ~ ., data = Boston, mtry = ncol(Boston) - 1, importance = TRUE)
Type of random forest: regression
Number of trees: 500
No. of variables tried at each split: 12
```

```
Mean of squared residuals: 11.5691
% Var explained: 84.95
```

**Question 2:** What is the *MSE* based on the training set?

**MSE = 11.5691**


How well does this bagged model perform on the test set?

**Question 3:** What is the formula to determine the *MSE*?

$$MSE = \frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2$$

Run the following in R.

```
yhat.bag = predict(bag.boston, newdata = Boston[-train,])
plot(yhat.bag, boston.test)
abline(0, 1)
```



```
mean((yhat.bag - boston.test)^2) ← Test MSE
```

```
[1] 23.23877
```

**Question 4:** What is the *MSE* of the test data set?

Test MSE = 23.23877

We could change the number of trees grown by `randomForest()` using the `ntree` argument:

```
bag.boston = randomForest(medv ~ ., data = Boston,
                           subset = train,
                           mtry = ncol(Boston) - 1,
                           ntree = 25)

bag.boston
```

Call:

```
randomForest(formula = medv ~ ., data = Boston, mtry = ncol(Boston) - 1, ntree = 25, s
              Type of random forest: regression
              Number of trees: 25
```

No. of variables tried at each split: 12

Mean of squared residuals: 12.30361

% Var explained: 83.99

```
yhat.bag = predict(bag.boston, newdata = Boston[-train,])  
mean((yhat.bag - boston.test)^2)
```

[1] 23.06258

**Question 5:** What method do we use to get the different trees?

We use bootstrap methods to get the different trees

## Random Forests

**Question 6:** For a building a random forest of regression trees, what should be `mtry` (number of predictors to consider at each split)?

$mtry = p/3$  ← Regression      $mtry = \sqrt{p}$  for classification

Type and run the following in R:

```
set.seed(10)  
rf.boston = randomForest(medv ~., data = Boston,  
                          subset = train,  
                          mtry = (ncol(Boston)-1)/3,  
                          importance = TRUE)  
yhat.rf = predict(rf.boston, newdata = Boston[-train,])  
mean((yhat.rf - boston.test)^2)
```

[1] 18.62328

**Question 7:** Compare the *MSE* of the test data to the *MSE* of the bagging.

This MSE is smaller than the results of bagging.

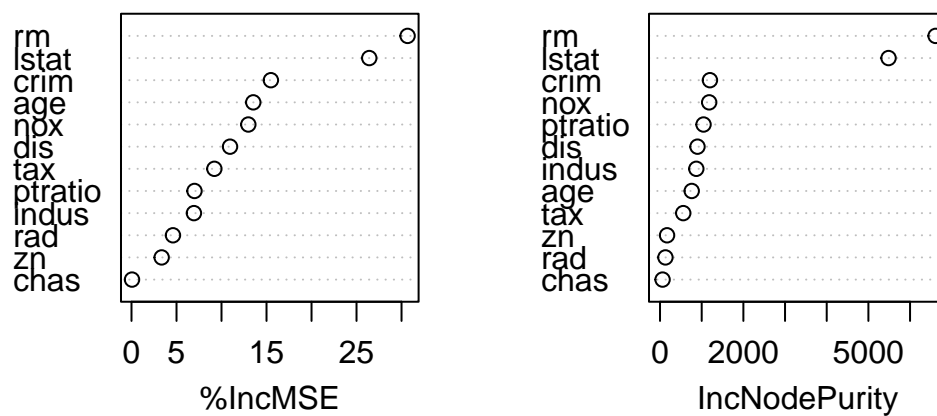
**Question 8:** Use the `importance()` function what are the two most important variables?

```
importance(rf.boston)
```

	%IncMSE	IncNodePurity
crim	15.48571304	1197.64717
zn	3.34978057	169.00931
indus	6.93488857	870.60348
chas	0.05746934	61.05778
nox	12.97835448	1179.66670
rm	30.67206810	6612.55554
age	13.52685213	760.41982
dis	10.94707995	899.17273
rad	4.60598124	129.80949
tax	9.20624202	556.89248
ptratio	6.99867017	1044.02812
lstat	26.41637352	5483.83696

```
varImpPlot(rf.boston)
```

rf.boston



rm and lstat


## Boosting

Run the following in R:

```

library(gbm)
set.seed(1)
boost.boston = gbm(medv ~., data = Boston[train,],
  distribution = "gaussian",
  n.trees = 5000,
  interaction.depth = 4)
summary(boost.boston)

```




	var	rel.inf
rm	rm	44.48249588
lstat	lstat	32.70281223
crim	crim	4.85109954
dis	dis	4.48693083
nox	nox	3.75222394
age	age	3.19769210
ptratio	ptratio	2.81354826
tax	tax	1.54417603
indus	indus	1.03384666
rad	rad	0.87625748
zn	zn	0.16220479
chas	chas	0.09671228

**Question 9:** What are the two most important variables with the boosted trees?


**rm and lstat**

We can produce *partial dependence plots* for these two variables. The plots illustrate the marginal effect of the selected variables on the response after *integrating* out the other variables.

```
plot(boost.boston, i = "rm")
```



```
plot(boost.boston, i = "lstat")
```



Notice that the house prices are increasing with `rm` and decreasing with `lstat`.

We will use the boosted model to predict `medv` on the test set:

```
yhat.boost = predict(boost.boston,
                      newdata = Boston[-train,],
                      n.trees = 5000)
mean((yhat.boost - boston.test)^2)
```

```
[1] 18.39057
```

**Question 10:** Compare this *MSE* to the *MSE* of the random forest and bagging models.

Bag MSE = 23.23877, RF MSE = 18.62328, Boost MSE = 18.39057