

Neural Networks Using Keras


MATH 4322

Instructions

- We will work on the `iris` data.
- We will use the `Keras` library to help build a neural network. Installation of the package is here <https://hastie.su.domains/ISLR2/keras-instructions.html>.
- You will also need to install the miniconda installation of python. This is where I went to install miniconda: <https://docs.conda.io/en/latest/miniconda.html>.
- Also you will need to install the package `corrplot` if you have not done so.


Exploratory Data Analysis

There are three species of iris flowers: versicolor, setosa and virginica.




For a comparison of the three iris species with respect to Petal.Length and Petal.Width, we can perform a scatterplot.

```
library(ggplot2)
ggplot(iris, aes(x = Petal.Length, y = Petal.Width, color = Species)) +
  geom_point()
```



Next we can look at the correlations of the variables with a correlation plot.

```
M = cor(iris[,1:4]) #Store the overall correlation matrix in M.
library(corrplot)
corrplot(M, method = "circle", type = "upper")
```



Separate into Training and Test subsets

Randomly divide the data into **training** and **testing** subset.

```
n <- nrow(iris)
p <- ncol(iris)-1

set.seed(1)
train <- sample(n, 0.8*n)

train_data <- iris[train, 1:p]
train_labels <- iris[train, p+1]

test_data <- iris[-train, 1:p]
test_labels <- iris[-train, p+1]
```

Random Forest

We will look at the random forest to predict the species. Then compare how well this works to the neural network.

```
library(randomForest)
set.seed(10)
rf.iris = randomForest(Species ~., data = iris,
                        subset = train,
                        mtry = sqrt(p),
                        importance = TRUE)
yhat.rf = predict(rf.iris, newdata = iris[-train,], type = "response")
tab = table(yhat.rf, iris[-train,]$Species)
```

Normalizing the data

For fit the neural network, it is best to scale the data so that each column has mean zero and variance one. This is called **Normalizing** the data.

```
train_data <- scale(train_data)
# Use means and standard deviations from training set to normalize test set
col_means_train <- attr(train_data, "scaled:center")
col_stddevs_train <- attr(train_data, "scaled:scale")

test_data <- scale(test_data,
                   center = col_means_train,
                   scale = col_stddevs_train)
```

One-Hot Encoding for the Categorical Reponse

When you want to model multi-class classification problems with the help of ANN, it is necessary to convert the response factor variable into numeric representation by [one-hot encoding \(OHE\)](#):

- our response `Species` can take on values `"setosa"`, `"versicolor"`, `"virginica"`
- We represent it via a **numerical** vector $\mathbf{x} = (x_1, x_2, x_3)$, such that

<i>Species</i>	x_1	x_2	x_3
"setosa"	1	0	0
"versicolor"	0	1	0
"virginica"	0	0	1

To do this, *keras* suggests function `to_categorical()`:

```
library(keras)
train_labels <- to_categorical(as.numeric(train_labels)-1)
head(train_labels)
```

```
      [,1] [,2] [,3]
[1,]    0    1    0
[2,]    0    0    1
[3,]    1    0    0
[4,]    1    0    0
[5,]    0    1    0
[6,]    0    1    0
```

```
test_labels <- to_categorical(as.numeric(test_labels)-1)
```

Next we want to set the seed for reproducibility.

```
tensorflow::tf$random$set_seed(1)
```

Building the Model with *Keras*

To start constructing a model, you should first [initialize a sequential model](#) with the help of the `keras_model_sequential()` function. We plan on constructing a network with

- a single hidden layer, which is fully-connected (`layer_dense`) and has 4 *ReLU* nodes (`units = 4`, `activation = "relu"`); the # of nodes for hidden layer is calculated via the rule of thumb

$$\# \text{ nodes in hidden layer} = \frac{\# \text{ nodes in input layer} + \# \text{ nodes in output layer}}{2} = \frac{4 + 3}{2} \approx 4$$

- a 3-node fully-connected output layer, which applies a `softmax` function.

```
library(embed) #https://github.com/tidymodels/embed
model <- keras_model_sequential(layers = list(
  layer_dense(units = 4, activation = 'relu', input_shape = dim(train_data)[2]),
  layer_dense(units = 3, activation = 'softmax')
))

model
```

Model: "sequential"

Layer (type)	Output Shape	Param #
dense (Dense)	(None, 4)	20
dense_1 (Dense)	(None, 3)	15

Total params: 35 (140.00 Byte)
 Trainable params: 35 (140.00 Byte)
 Non-trainable params: 0 (0.00 Byte)

Compile and Train the Model

We need to compile the model by specifying:

- **loss function**, *loss* - what criteria does our ANN optimize?
- **optimizing algorithm**, *optimizer* - how do we optimize (details left out),
- **model performance metrics**, *metrics*.

```
compile(model,
  loss = 'categorical_crossentropy',
  optimizer = 'adam',
  metrics = 'accuracy'
)
```

Since we have a classification problem, here the metric is the overall prediction accuracy which is:

$$accuracy = \frac{\# \text{ of correct predictions}}{\# \text{ of all predictions}}$$

Next, we proceed to `train`, or `fit`, the model while using our `training subset`. We will train it

- for `200 epochs`,
- feeding `batch_size = 32` input samples at a time,
- holding out `20%` of the `training data` to be used for `out-of-sample validation` of the training process. This validation set will provide us with ability to track `model's accuracy performance` on `data that wasn't used for training`.

```
# Store the fitting history in `history`  
history <- fit(model,  
               train_data, train_labels,  
               epochs = 200,  
               batch_size = 32,  
               validation_split = 0.2)
```


When `verbose = 0` option is `not specified`, the progress printouts will be showing up after each epoch. Those will include:

- the `loss` and `accuracy` values on `96 training samples` (out of 120 total training samples),
- the `loss` and `accuracy` values on `24 validation samples` (out of 120 training samples).

```
history
```

We can also plot this:

```
plot(history, smooth=F)
```



UPDATE: MORE EPOCHS and EARLY STOPPING

To get better accuracy we could use more epochs. Here we will use 1000 epochs and prevent overfitting by early stopping after 20 epochs of no improvement to the training cross-entropy loss.

```
#use_session_with_seed(1)
tensorflow::tf$random$set_seed(1)

## Early stopping callback function
early_stop <- callback_early_stopping(monitor = "val_loss",
                                       patience = 20)

#### Initializing & Building the model with KERAS

# Initialize a sequential model
model <- keras_model_sequential(layers = list(
  layer_dense(units = 4, activation = 'relu', input_shape = dim(train_data)[2]),
  layer_dense(units = 3, activation = 'softmax')
))

model
```



```
### COMPILER the model.


# Compiling the model

compile(model,
        loss = 'categorical_crossentropy',
        optimizer = 'adam',
        metrics = 'accuracy'
)

### FIT THE MODEL

# Store the fitting history in `history`
history <- fit(model,
              train_data, train_labels,
              epochs = 1000,
              batch_size = 32,
              validation_split = 0.2,
              callbacks=list(early_stop))

# Plot the history
plot(history, smooth=F)
```



```
# Print the training loss and metrics
history
```

Test Data Predictions

We can use the model to predict the species for the test subset.

```
test.pred = predict(model, test_data)
head(test.pred, 2)
tail(test.pred, 2)
```

This provides the calculated class probabilities for each species.

We can put these predictions into a confusion matrix.

```
test_labels_num <- as.numeric(iris[-train, p+1] )-1
model %>% predict(test_data) %>% k_argmax()

# Evaluate on test data and labels
evaluate(model,
```



```
test_data, test_labels)
```